

Assignment 3 - Sparse Matrix Vector Multiplication

Kristian Sørdal

May 7, 2024

Contents

1	Sparse Matrix Vector Multiplication (SpMV)	3
1.1	Representing a Sparse Matrix effectively	3
1.2	SpMV Kernel	4
2	STREAM Benchmark	4
3	Parallelizing SpMV	5
3.1	Method 1	5
3.2	Method 2	5
3.3	Method 3	6
3.4	Load Balance comparison	8

1.2 SpMV Kernel

The kernel used for computing the result of an iteration of SpMV on a CSR matrix is outlined in the algorithm below. This is what will be parallelized in order to (hopefully) achieve better performance.

Algorithm 1: SpMV Kernel

Input : Row pointer: $rowPtr$, Column indices: $colIdx$, Values: $vals$, Input vector x , Result vector y

Output: Output Vector y

```

for  $row \leftarrow 0$  to  $n$  do
  for  $idx \leftarrow rowPtr[row]$  to  $rowPtr[row + 1]$  do
     $y[row] \leftarrow y[row] + vals[idx] \times x[colIdx[idx]]$ 

```

2 STREAM Benchmark

When performing SpMV with the CSR format, we have to be aware that the performance of our program will likely not scale with the number of threads we throw at the program, but rather the memory bandwidth of the system we are running on. Memory bandwidth is the rate at which data can be read or stored in memory. The reason for this becomes evident when we look at the kernel.

When we execute the kernel, we are repeatedly accessing memory, and do very few operations (just 2 floating point operations) per memory access. In addition to this, we are accessing the input vector x in a manner that leads to a large amount of cache misses. This is because it is indexed by the value stored in col_idx . Although this vector is sorted in increasing order when looking at a single row, the indices stored in this vector can, and usually will have a large difference in their numerical value. When we access any value at a given index in the x vector, the processor will prefetch successive elements in this vector and load them into the cache, allowing for very fast access. But since the numerical difference between the indices can be large, most likely the next element we need will not be prefetched, thus requiring the processor to read the value from memory, instead of reading from the cache. It is for this reason that SpMV is limited by the memory bandwidth of the system.

The STREAM is an industry standard benchmark that is used for measuring sustained memory bandwidth of the CPU in a shared memory system. The benchmark computes four different kernels named **Copy**, **Scale**, **Sum** and **Triad**. The kernels are defined as the following:

- Copy: $y[i] = x[i]$
- Scale: $y[i] = \alpha \times x[i]$
- Sum: $y[i] = x[i] + y[i]$
- Triad: $y[i] = x[i] + \alpha \times z[i]$

We are interested in the results from the **Triad** kernel, as this is the only kernel which uses 2 FLOPS per iteration, which is the same as we do in SpMV.

By running the STREAM benchmark on Brake and plotting the results, we can see the sustained memory bandwidth achieved for each thread. This will be important for comparing the results of the different implementations of parallel SpMV later on. Below we can see the results of the benchmark.

As we can see, we are able to achieve at most ≈ 33 GB/s in sustained memory bandwidth, and a peak of ≈ 2.5 GFLOPS, so we should not expect that the SpMV results will exceed this number.

3 Parallelizing SpMV

For this project, 3 different strategies were implemented for parallelizing SpMV. The following sections will discuss their implementation details.

3.1 Method 1

This method is the simplest of all methods, as it simply parallelizes the outmost for-loop of the kernel, using OpenMP's `#pragma omp parallel for` directive. From testing different scheduling methods, it was found that `dynamic` scheduling with a block size of 1024 was the best choice. The kernel now looks like the following:

Algorithm 2: SpMV Kernel

Input : Row pointer: *rowPtr*, Column indices: *colIdx*, Values: *vals*, Input vector *x*, Result vector *y*

Output: Output Vector *y*

```

for row  $\leftarrow$  0 to n do in parallel
    for idx  $\leftarrow$  rowPtr[row] to rowPtr[row + 1] do
         $y[\text{row}] \leftarrow y[\text{row}] + \text{vals}[\text{idx}] \times x[\text{colIdx}[\text{idx}]]$ 

```

3.2 Method 2

The next method is slightly more involved, as it makes use of manually assigning rows to each thread, in an attempt at achieving a better load balance for each thread, trying to ensure that no thread does significantly more work than any of the other threads.

The load balancing strategy used in this method involves computing the ideal number of non-zeros each thread should work on, by dividing the total number of non-zeros by the number of threads. Then for each thread, we greedily assign rows (without splitting them) until the sum of non-zeros in each row is at or just above the ideal number of non-zeros. The algorithm used for this is implemented in the following manner:

Algorithm 3: Naive Load Balancing

Input : Row pointer: *rowPtr*, Number of non-zeros: *nnz*, Number of threads: *numThreads*, Number of rows: *N*,

Output: Start and end indices for each rank
 $\text{partitionSize} \leftarrow$ list of size *p*, assigned tuples (0, 0)

$\text{avgNnz} \leftarrow \frac{\text{nnz}}{p}$

$\text{threadIdx} \leftarrow 0$

$\text{prevIdx} \leftarrow 0$

for *i* \leftarrow 0 **to** *N* **do**

if $\text{threadIdx} = \text{numThreads} - 1$ **then**
 $\text{partition}[\text{threadIdx}] \leftarrow (\text{prevIdx}, N)$ **break**

if $\text{rowPtr}[i+1] - \text{rowPtr}[\text{prevIdx}] \geq \text{avgNnz}$ **then**
 $\text{partition}[\text{threadIdx}] \leftarrow (\text{prevIdx}, i)$ $\text{prevIdx} \leftarrow i$ $\text{threadIdx} \leftarrow \text{threadIdx} + 1$

Then, for each thread, we compute the following kernel:

Algorithm 4: SpMV Kernel - Load Balanced

Input : Row pointer: $rowPtr$, Column indices: $colIdx$, Values: $vals$, Input vector x , Result vector y , Start index: $startIdx$, End index: $endIdx$

Output: Output Vector y

```
for  $row \leftarrow startIdx$  to  $endIdx$  do
  for  $idx \leftarrow rowPtr[row]$  to  $rowPtr[row + 1]$  do
     $y[row] \leftarrow y[row] + vals[idx] \times x[colIdx[idx]]$ 
```

3.3 Method 3

The final method implemented makes use of the METIS graph partitioning library, this exposes two methods for partitioning a graph into k parts. We have the choice of using either multilevel recursive bisection, or k -way partitioning. After some testing, k -way partitioning was found to be the most stable, as the multilevel recursive bisection sometimes resulted in segmentation faults.

By using the METIS library, the goal is to hopefully achieve an even better load balance than what the greedy approach described earlier can provide. The algorithm used for load balancing with METIS is outlined below.

Algorithm 5: Metis Load Balancing

Input : Number of partitions: k , Row pointer: $rowPtr$, Column indices: $colIdx$, Values: $vals$, Input vector: $inputVec$, Number of non-zeros: nnz , Number of rows: N

Output: Start and end indices for each rank

```
startIdx  $\leftarrow$  0 initialized list of size  $p + 1$ 
partitionMap  $\leftarrow$  0 initialized list of size  $N$ 
objVal  $\leftarrow$  NULL
ncon  $\leftarrow$  1
```

```
//ubvec constrains the load imbalance, a value of 1.01 indicates that the load
  imbalance is at most 1%. Note that there is no guarantee that this is
  achieved
ubvec  $\leftarrow$  1.01
```

```
if  $k = 1$  then
  startIdx[1] = N return
```

```
//METIS partitioning
```

```
rc  $\leftarrow$  METISPartGraphKway(&N, &ncon, *rowPtr, *colIdx, nullptr, nullptr, nullptr, &k,
  nullptr, &ubvec, nullptr, &objval, *partitionMap)
```

```
newId  $\leftarrow$  0 initialized list of size  $N$ 
oldId  $\leftarrow$  0 initialized list of size  $N$ 
id  $\leftarrow$  0
//Reenumerate the rows according to the partitionMap
for  $r \leftarrow 0$  to  $k$  do
  for  $i \leftarrow 0$  to  $N$  do
    if partitionMap[i] =  $r$  then
      newId[id]  $\leftarrow$  i oldId[i]  $\leftarrow$  id id  $\leftarrow$  id + 1
  startIdx[r+1]  $\leftarrow$  id
```

```
newRowPtr  $\leftarrow$  0 initialized list of size  $N + 1$ 
newInputVec  $\leftarrow$  0 initialized list of size  $N + 1$ 
newColIdx  $\leftarrow$  0 initialized list of size  $nnz$ 
newVals  $\leftarrow$  0 initialized list of size  $nnz$ 
```

```
for  $i \leftarrow 0$  to  $N$  do
  degree  $\leftarrow$  rowPtr[i+1] - rowPtr[i]
  newRowPtr[i+1] = newRowPtr[i] + degree
  colIdxStart  $\leftarrow$  *colIdx + rowPtr[oldId[i]]
  valsStart  $\leftarrow$  *vals + rowPtr[oldId[i]]

  //Copy the values and indices to the new arrays
  copy(colIdxStart, colIdxStart + degree, newColIdx + newRowPtr[i])
  copy(valsStart, valsStart + degree, newVals + newRowPtr[i])

  for  $r \leftarrow$  newRowPtr[i] to newRowPtr[i+1] do
    newColIdx[r]  $\leftarrow$  newId[newColIdx[r]]
```

```
for  $i \leftarrow 0$  to  $N$  do
  newInputVec[i] = inputVec[oldId[i]]
for  $i \leftarrow 0$  to  $k$  do
  partition[i]  $\leftarrow$  (startIndices[i], startIndices[i+1])
rowPtr  $\leftarrow$  newRowPtr
colIdx  $\leftarrow$  newColIdx
vals  $\leftarrow$  newVals
inputVec  $\leftarrow$  newInputVec
```

3.4 Load Balance comparison